

Measuring Cross-Lingual Robustness for Chemistry-Patent Retrieval: Patent-Grounded Multilingual Benchmarks, New Metrics, and a Deployment Decision

Anonymous ACL submission

Abstract

Chemical knowledge is produced, searched, and reused across language boundaries: a technical question may be asked in one language while the strongest evidence appears in another. Yet multilingual chemistry retrieval has little clean public evaluation data, so models are often judged by aggregate recall even when that number hides the failure that matters most: retrieving the right evidence across languages rather than an easier same-language match. We address this gap with a multilingual chemistry QAC benchmark and generation pipeline for cross-lingual retrieval evaluation. Patents provide the benchmark substrate in this work because multilingual publication variants and chemistry ontologies make it possible to construct controlled relevance judgments and chemically plausible hard negatives. The pipeline extracts chemistry-relevant technical text, generates multiple question-answer-context candidates with an LLM, grades them for faithfulness and retrieval quality, and validates the automated grading against human annotation. We export the resulting data in retrieval-ready form with corpus, queries, qrels, and QAC records, enabling evaluation by query language, document language, same-language versus cross-language relevance, and chemically confusable hard negatives. Across evaluated multilingual embedding models, average performance masks systematic cross-language weakness: models can prefer same-language or chemically similar but irrelevant documents over the intended evidence. The benchmark and pipeline provide a practical framework for building and auditing specialized multilingual retrieval evaluations in high-trust technical domains. Data and code are available on [Hugging Face](#) and [GitHub](#).

1 Introduction

Technical search in chemistry rarely respects a single language boundary. A researcher may ask a question in English, German, French, Spanish, or

another working language, while the strongest evidence may be documented elsewhere: in a translated patent family, in a regional filing, or in technical text whose terminology follows a different linguistic convention. This is not only a translation problem. Chemistry search depends on aliases, abbreviations, formulae, compound families, process descriptions, and near-neighbor concepts that can be semantically close while still being wrong for the user’s question. For a retrieval system used in a high-trust technical workflow, returning a plausible but chemically different document is not a harmless ranking error.

Dense retrieval with neural sentence and passage embeddings (Reimers and Gurevych, 2019; Karpukhin et al., 2020) is attractive for this setting because multilingual embedding models promise a single retrieval index over many languages (Feng et al., 2022; Chen et al., 2024). In practice, however, the evaluation signal used to choose such a model is often too coarse. A model can achieve a strong average Recall@k by retrieving easy same-language evidence, while failing when the intended evidence is in another language, a bias also observed in recent cross-lingual retrieval analyses (Goworek et al., 2025; Amiraz et al., 2025). The problem is not only whether a multilingual encoder retrieves a relevant document, but whether it retrieves the right document when query language, document language, and chemical similarity pull the ranking in different directions. This is the failure mode that matters in deployment, and it is precisely the one that an aggregate retrieval score can hide.

Existing multilingual retrieval benchmarks provide important coverage for general cross-lingual search (Zhang et al., 2023; Lawrie et al., 2023; Muennighoff et al., 2023; Enevoldsen et al., 2025), but they do not directly capture this chemistry workflow. We need evaluation data where language variants are controlled, questions are

grounded in technical contexts, and hard negatives are chemically plausible rather than arbitrary. Without this structure, an evaluation may reward the model that looks best on average while missing the model behavior that creates practical risk.

We present a multilingual chemistry question-answer-context (QAC) benchmark and generation pipeline for cross-lingual retrieval evaluation. The broader task is multilingual chemistry retrieval; patents are the substrate we use to construct the benchmark because they provide multilingual technical variants and chemistry-rich descriptions at scale, building on the long history of patent retrieval evaluation (Piroi et al., 2011). The benchmark is designed to make hidden retrieval failures visible: it separates same-language and cross-language relevance, supports analysis by query and document language, and includes chemistry-aware hard negatives derived from alias and ontology structure, including ChEBI (Hastings et al., 2016).

This paper is therefore an evaluation-framework paper rather than a model leaderboard. The contribution is a practical way to build, validate, and audit specialized multilingual retrieval benchmarks in a domain where both language and technical similarity matter. Our contributions are a public multilingual chemistry QAC benchmark, a reproducible LLM-assisted generation and validation pipeline, a language-aware retrieval evaluation setup, and a chemistry-confusability stress test based on alias and ontology structure.

Together, these components provide a reusable lesson for high-trust technical retrieval: start from the workflow failure, build data that makes the failure observable, and report language-aware robustness alongside standard retrieval metrics.

2 Related Work

Multilingual retrieval benchmarks such as MIRACL, NeuCLIR, MTEB, MMTEB, and CLIRMatrix have made cross-lingual retrieval evaluation easier to compare across models and languages (Zhang et al., 2023; Lawrie et al., 2023; Muenighoff et al., 2023; Enevoldsen et al., 2025; Sun and Duh, 2020). They are closest to our evaluation setting, but they are not designed around chemistry-specific relevance, controlled language variants, or chemically plausible hard negatives. Our benchmark follows this line of work while adding domain structure needed to audit multilin-

gual retrieval in technical chemistry search.

Recent work on cross-lingual retrieval and multilingual RAG shows that model rankings can depend strongly on language, alignment, and same-language preference (Goworek et al., 2025; Amiraz et al., 2025; Li et al., 2024; Liu et al., 2025; Ki et al., 2025). These findings motivate our focus on separating same-language from cross-language relevance instead of relying only on aggregate recall. In contrast to settings where content and language are difficult to disentangle, patent publication variants give us a practical way to compare multilingual technical evidence under tighter content control.

Patent retrieval has a long evaluation history, including NTCIR and CLEF-IP (Fujii and Ishikawa, 2002; Piroi et al., 2011), and recent work has revisited patent embeddings (Ghosh et al., 2024; Yousefiramandi and Cooney, 2026). Our goal is complementary: we do not propose a new encoder, but a chemistry-grounded evaluation framework. We use chemical ontology and alias structure, including ChEBI (Hastings et al., 2016), to define hard negatives that are plausible for chemistry search but wrong for the intended evidence.

3 QAC Generation Pipeline

Our method turns multilingual chemistry technical text into auditable question-answer-context (QAC) records for retrieval evaluation. The design goal is not simply to generate fluent questions. Each generated item must identify the evidence document it came from, preserve language metadata, and carry quality signals that make the benchmark inspectable after export. Figure 1 summarizes the full pipeline.

Corpus and context construction. The input is a chemistry-relevant multilingual corpus. In this work, we use patent-derived text as the benchmark substrate because patents provide technical descriptions and multilingual publication variants at scale. For each document-language pair, the context builder assembles a compact evidence block from available fields such as title, abstract, and technical description or claim text. This step keeps the generation grounded in a specific document rather than asking the model to produce generic chemistry questions from memory.

Candidate generation. For each selected document, language, and generation mode, an LLM




Figure 1: Overview of the multilingual chemistry retrieval benchmark pipeline, from chemistry technical text through QAC generation, verification, human validation, benchmark export, and retrieval evaluation.

Figure note: Icons adapted from Font Awesome Free, licensed under CC BY 4.0.

produces multiple QAC candidates. Each candidate contains a user-facing question, a short answer, the supporting context, and metadata linking the item back to the source document and language. We use both technical and semantic question modes so that the benchmark contains queries that reflect different retrieval intents: some ask for concrete technical details, while others ask for a higher-level description of what the document is about.

Verification and filtering. Generated candidates are then passed through an LLM-based verifier. The verifier scores whether the answer is faithful to the provided context, whether the question is answerable from that context, whether the language is appropriate, and whether the item is useful as a retrieval query. We retain the best-scored items together with their audit metadata rather than treating generation as a black box. This gives each QAC row a provenance trail: source document, language, generation mode, scores, and verifier feedback.

Human validation and export. Because the verifier is itself model-based, we validate its behavior against a human-annotated sample. This does not replace expert review for every item, but it checks whether the automatic grading is directionally reliable enough for benchmark construction. The filtered and validated QAC records are then exported with the retrieval artifacts used in later sections: a corpus table, query table, rele-

vance judgments, and the full QAC table. These artifacts let us evaluate models with standard dense-retrieval tooling while still slicing results by query language, document language, and source provenance.

4 Benchmark Construction

The filtered QAC records are exported as retrieval benchmarks rather than as a single flat table. Each release contains four linked views: corpus stores the searchable documents and their languages; queries stores the generated questions; qrels stores relevance judgments from queries to documents; and qac preserves the original question, answer, context, source document, generation metadata, and quality scores. This separation lets standard retrieval evaluators consume the data while keeping the provenance needed to audit failures by language and source. Table 1 summarizes the statistics we rely on in this short paper.

Cross-language relevance. For the multilingual chemistry retrieval benchmark, each query is tied to a source technical document. Relevant documents include the source document and its available multilingual variants, so a query can be evaluated against both same-language and cross-language evidence. We keep query language, document language, and source identifiers as explicit metadata. This makes it possible to separate an easy same-language hit from a harder cross-language hit, and to construct “no-home” cases in

Item	Value	
Chem-patents corpus	rows	1.11k documents
Chem-patents queries	QAC /	137 QAC rows / 137 queries
Chem-patents qrels		322 total; 265 cross-language qrels
Languages		5 (English, German, Spanish, French, Chinese)
Alias-graph questions		132 questions over 24 ChEBI concepts
Alias-graph hard negatives		Chemically confusable ontology / alias neighbors
Human validation sample		97 reviewed QACs

Table 1: Compact benchmark statistics used in this draft. Chem-patents counts come from the released Hugging Face dataset card; alias-graph counts come from the benchmark analysis used by the full draft.

which the query language has no same-language gold document. These cases are central to the benchmark because they remove the shortcut that often inflates aggregate recall.

Chemistry-aware hard negatives. Language is only one source of retrieval error in chemistry. A model may also rank a document about a related compound, synonym, parent class, or neighboring concept above the intended evidence. To test this failure mode, we build an alias-graph stress test from chemistry ontology and alias structure. ChEBI concepts and multilingual aliases define target chemical concepts; documents associated with the target concept form the gold set; and taxonomic or alias-graph neighbors provide chemically plausible hard negatives. These negatives are deliberately more informative than random non-relevant documents: they ask whether the retriever can distinguish the right chemical evidence from evidence that is close enough to be tempting.

The resulting benchmark labels each retrieved document by its role in the workflow: same-language relevant evidence, cross-language relevant evidence, unrelated distraction, or chemically confusable hard negative. The next section uses these labels to evaluate multilingual embedding models beyond aggregate recall.

5 Evaluation Setup

We evaluate multilingual embedding models as drop-in dense retrievers. Each model embeds the query and all corpus documents; documents are ranked by embedding similarity; and retrieval quality is computed from the benchmark qrels.




Figure 2: Evaluation separates aggregate Recall@10 from same-language retrieval and cross-language retrieval. The split is the central diagnostic: a model can look strong on average while still losing substantial ground when the relevant evidence is in another language.

This setup matches the practical decision faced by a technical search team: choose one multilingual encoder for a shared index, then ask whether it retrieves the right evidence across languages.

Models and protocol. We evaluate nine multilingual or domain-relevant embedding models: embeddinggemma, bge-m3 (Chen et al., 2024), granite-278m, nomic-v2-moe, qwen3-0.6B, LaBSE (Feng et al., 2022), SapBERT (Liu et al., 2021), e5-large-instruct, and gte-base. All models retrieve against the same multilingual chemistry corpus, and all queries are evaluated with the same relevance judgments. We use standard Recall@k as the baseline retrieval metric because it is familiar and easy to compare, but we treat it as only the starting point: average recall does not say where the relevant document came from or which failure mode produced a miss.

Language-aware views. For each model, we split results by query language, document language, and whether a retrieved relevant document is in the same language as the query or in another language. This gives two central views: same-language retrieval, where a model can often exploit lexical or embedding alignment advantages, and cross-language retrieval, where the model must bridge the language boundary. We also inspect language-pair behavior as a directed matrix, since retrieving a German document from an English query need not be as easy as retrieving an English document from a German query.

Failure analysis. Beyond relevant-document recall, we measure the kinds of mistakes that matter for the workflow. First, we track same-language distractions: cases where a model retrieves a non-




Figure 3: The alias-graph stress test measures chemistry-confusability: cases where a chemically plausible but incorrect document outranks all gold evidence. This complements language-aware recall by testing whether models distinguish the intended chemical concept from near neighbors.

relevant document in the query language ahead of relevant cross-language evidence. Second, we use the alias-graph benchmark to measure chemistry confusion, where a document about a chemically related but incorrect concept outranks the intended evidence. These analyses turn retrieval output into diagnostic categories rather than a single score.

Reporting principle. The results therefore report standard retrieval metrics together with language-aware and chemistry-aware slices. A model is useful for this setting only if it performs well on the aggregate task and remains robust when the relevant evidence changes language or when chemically plausible negatives enter the ranking.

6 Results

We organize the results around the three questions the benchmark is meant to answer: whether generated QACs are usable, whether cross-language retrieval is harder than same-language retrieval, and whether chemistry-aware hard negatives surface failures that ordinary relevance metrics hide.

QAC quality is high enough for benchmark construction. Human review of 97 generated QACs gives a mean score of 8.33/10, with 94 of 97 items falling in the “good” bucket. Technical questions score slightly higher than semantic questions on average (8.72 vs. 7.91), but both modes remain usable for retrieval evaluation. Comparing normalized human and LLM-grader scores, the LLM grader is slightly stricter than the human annotator on average (79.0% vs. 83.3%), with a mean absolute difference of 8.22 percentage points. We read this as evidence that the automatic verifier is




Figure 4: Language-collapse diagnostics for the cross-lingual benchmark. Same-language documents are over-represented in retrieved lists, and same-language distractors can bury cross-language gold evidence. This explains why average recall can hide the failure most relevant to multilingual chemistry search.

useful for filtering and auditing, while still requiring human validation for claims about benchmark quality.

Average recall hides cross-language weakness. The cross-lingual benchmark shows that standard Recall@10 is not enough to judge a multilingual retriever. In Figure 2, same-language retrieval is consistently easier than cross-language retrieval: the best cross-language Recall@10 is 0.50, and the same-language home advantage reaches +0.55 for the most biased model. This means a model can look acceptable under an aggregate score while failing on the deployment-relevant case where the evidence is available only in another language. The directed language-pair view further shows that the difficulty is not symmetric: query language and document language both matter.

The errors are not only multilingual; they are chemical. The alias-graph benchmark shows a second failure mode. Even when a retrieved document is topically close, it may be chemically wrong. Figure 3 measures how often a chemically confusable document outranks all gold evidence; on the publication lens, this happens on 14–78% of queries depending on the model. These failures are especially important for chemistry search because a nearby compound, parent class, or sibling concept can look plausible while answering the wrong question.

Together, the results support the central lesson of the benchmark: multilingual chemistry retrieval should be evaluated with language position and chemical confusability visible, not only with a single averaged retrieval score.

7 Deployment Decision

The evaluation is meant to support a practical model-selection decision: choose one multilingual




Figure 5: Deployment-oriented cost-vs-capability frontier. Cross-language Recall@10 measures capability, while XRC50 measures how much deeper a user must read to reach cross-language evidence. embeddinggemma is the maximum-capability corner; bge-m3 is the cheaper-to-read option at a CLIR@10 threshold of 0.40.

embedding model for a shared chemistry retrieval index. If the team looked only at aggregate retrieval scores, the decision would be hard to audit because same-language hits and chemically plausible distractors are mixed into one number. Our benchmark changes the decision criterion from “which model has the best average Recall@10?” to “which model remains usable when the evidence is cross-language and chemically confusable?”

Capability vs. reading cost. Figure 5 shows the resulting cost-vs-capability view. We use cross-language Recall@10 as the capability axis and XRC50, the median cross-lingual reading-cost multiplier, as the cost axis. The Pareto frontier has three models: embeddinggemma, bge-m3, and granite-278m. embeddinggemma is the unique maximum-CLIR corner (CLIR@10 = 0.5024, XRC50 = 3.5), so it is the recommended capability-first choice. If an operating threshold of $\text{CLIR@10} \geq 0.40$ is acceptable, bge-m3 becomes the cheaper-to-read admitted model (XRC50 = 2.0). At a lower threshold, granite-278m enters as an even cheaper frontier point (XRC50 = 1.25, CLIR@10 = 0.3285). Thus the benchmark does not produce a context-free winner; it exposes the trade-off a deployment team must choose.

Decision rule. For a deployment that prioritizes maximum cross-language recall, we would select embeddinggemma. For a deployment with a stricter reading-cost budget and an acceptable CLIR@10

threshold around 0.40, bge-m3 is the more efficient alternative. In both cases, the decision uses language-aware and chemistry-aware diagnostics rather than aggregate recall alone. This is the practical lesson: model selection for high-trust multilingual chemistry search should report capability, language robustness, and chemistry-confusability side by side before a retriever is shipped.

Limitations

This benchmark is intentionally specialized. Chemistry patents provide a useful substrate for controlled multilingual technical evidence, but results may not transfer directly to other scientific domains, regulatory corpora, or general web retrieval. The design should be read as a reusable evaluation pattern rather than as proof that the same model ordering will hold everywhere.

Patent publication variants are also an imperfect proxy for equivalent technical evidence. They give stronger content control than unrelated multilingual corpora, but they do not guarantee full claim-level equivalence across every language version. We therefore use them as controlled evidence for retrieval evaluation, not as a legal statement that all patent variants are interchangeable.

The QAC pipeline relies on LLM generation and LLM-based verification. Human review supports the usefulness of the generated data, but the validation sample is limited and does not replace expert patent or chemistry review for every item. Future versions should expand expert annotation, especially for borderline chemical relevance and hard-negative severity.

Finally, the current paper focuses on multilingual retrieval and alias-graph confusability. Code-switching and noisy-query infrastructure exist in the project, but we do not claim complete results for those settings here.

8 Conclusion

Multilingual chemistry retrieval needs evaluation data that makes the right failure modes visible. In this paper, we presented a QAC generation and validation pipeline for building such data, released it in retrieval-ready form, and used it to evaluate dense multilingual retrievers beyond aggregate recall. The central lesson is simple: a model that looks effective on average may still prefer same-language shortcuts or chemically plausible but wrong documents over the intended cross-

language evidence.

The benchmark shows how to make those failures measurable. Language metadata separates same-language from cross-language relevance; QAC provenance makes generated queries auditable; human validation checks the LLM-based verifier; and alias-graph hard negatives expose chemistry-specific confusions. These pieces are useful beyond the particular patent-derived corpus used here: they provide a template for building specialized multilingual retrieval evaluations in high-trust technical domains.

For practitioners, the recommendation is to report robustness next to recall. For benchmark builders, the recommendation is to start from the workflow failure and preserve enough metadata to diagnose it. In multilingual chemistry search, that means asking not only whether a relevant document was retrieved, but whether the system found the intended evidence across language and chemical confusability boundaries.

References

- Chen Amiraz, Yaroslav Fyodorov, Elad Haramaty, Zohar Karnin, and Liane Lewin-Eytan. 2025. [The cross-lingual cost: Retrieval biases in RAG over Arabic-English corpora](#). *arXiv preprint arXiv:2507.07543*.
- Jianlv Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. 2024. [BGE M3-Embedding: Multi-lingual, multi-functionality, multi-granularity text embeddings through self-knowledge distillation](#). *arXiv preprint arXiv:2402.03216*.
- Kenneth Enevoldsen and 1 others. 2025. [MMTEB: Massive multilingual text embedding benchmark](#). *arXiv preprint arXiv:2502.13595*.
- Fangxiaoyu Feng, Yinfei Yang, Daniel Cer, Naveen Arivazhagan, and Wei Wang. 2022. Language-agnostic BERT sentence embedding. *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 878–891.
- Atsushi Fujii and Tetsuya Ishikawa. 2002. [NTCIR-3 patent retrieval experiments at ULIS](#). *arXiv preprint cs/0206035*.
- Mainak Ghosh, Sebastian Erhardt, Michael E. Rose, Erik Buunk, and Dietmar Harhoff. 2024. [PaECTER: Patent-level representation learning using citation-informed transformers](#). *arXiv preprint arXiv:2402.19411*.
- Roksana Goworek, Olivia Macmillan-Scott, and Eda B. Özyiğit. 2025. [What drives cross-lingual ranking? retrieval approaches with multilingual language models](#). *arXiv preprint arXiv:2511.19324*.
- Janna Hastings, Gareth Owen, Adriano Dekker, Marcus Ennis, Namrata Kale, Venkatesh Muthukrishnan, Steve Turner, Neil Swainston, Pedro Mendes, and Christoph Steinbeck. 2016. ChEBI in 2016: Improved services and an expanding collection of metabolites. *Nucleic Acids Research*, 44(D1):D1214–D1219.
- Vladimir Karpukhin, Barlas Öguz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen tau Yih. 2020. [Dense passage retrieval for open-domain question answering](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 6769–6781.
- Dayeon Ki, Marine Carpuat, Paul McNamee, Daniel Khashabi, Eugene Yang, Dawn Lawrie, and Kevin Duh. 2025. [Linguistic nepotism: Trading-off quality for language preference in multilingual RAG](#). *arXiv preprint arXiv:2509.13930*.
- Dawn Lawrie, Sean MacAvaney, James Mayfield, Paul McNamee, Douglas W. Oard, Luca Soldaini, and Eugene Yang. 2023. Overview of the TREC 2022 NeuCLIR track. In *Proceedings of the Thirty-First Text REtrieval Conference (TREC)*.
- Bryan Li, Fiona Luo, Samar Haider, Adwait Agashe, Tammy Li, Runqi Liu, Muqing Miao, Shriya Ramakrishnan, Yuan Yuan, and Chris Callison-Burch. 2024. [BordIRlines: A dataset for evaluating cross-lingual retrieval-augmented generation](#). *arXiv preprint arXiv:2410.01171*.
- Fangyu Liu, Ehsan Shareghi, Zaiqiao Meng, Marco Basaldella, and Nigel Collier. 2021. Self-alignment pretraining for biomedical entity representations. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pages 4228–4238.
- Wei Liu, Sony Trenous, Leonardo F. R. Ribeiro, Bill Byrne, and Felix Hieber. 2025. [XRAG: Cross-lingual retrieval-augmented generation](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 15669–15690.
- Niklas Muennighoff, Nouamane Tazi, Loïc Magne, and Nils Reimers. 2023. MTEB: Massive text embedding benchmark. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, pages 2014–2037.
- Florina Piroi, Mihai Lupu, Allan Hanbury, and Veronika Zenz. 2011. [CLEF-IP 2011: Retrieval in the intellectual property domain](#). In *CLEF (Notebook Papers/Labs/Workshop)*. Multilingual (en/de/fr) patent prior-art retrieval; see also CLEF-IP 2009–2013 overviews.
- Nils Reimers and Iryna Gurevych. 2019. [Sentence-BERT: Sentence embeddings using Siamese BERT-networks](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3982–3992.
- Shuo Sun and Kevin Duh. 2020. [CLIRMatrix: A massively large collection of bilingual and multilingual datasets for cross-lingual information retrieval](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 4160–4170.
- Amirhossein Yousefiramandi and Ciaran Cooney. 2026. [Benchmarking patent embeddings: A multi-task evaluation of 22 models across retrieval, classification, and clustering](#). *arXiv preprint arXiv:2605.24297*.
- Xinyu Zhang, Nandan Thakur, Odunayo Ogundepo, Ehsan Kamalloo, David Alfonso-Hermelo, Xiaoguang Li, Qun Liu, Mehdi Rezagholizadeh, and Jimmy Lin. 2023. [MIRACL: A multilingual retrieval dataset covering 18 diverse languages](#). *Transactions of the Association for Computational Linguistics*, 11:1114–1131.




Figure 6: Same-language home advantage by model. The gap between same-language and cross-language retrieval is the main reason aggregate recall can overstate deployment readiness.




Figure 7: Directional query-language $\rightarrow$ document-language retrieval. Cross-language retrieval is not symmetric: a model’s behavior depends on both the query language and the target document language.

A Appendix

This appendix collects supporting analyses that are useful for interpretation but too detailed for the six-page body.

B Additional Cross-Lingual Retrieval Diagnostics

C Additional Alias-Graph Diagnostics

D Deployment-Oriented Supporting Analyses




Figure 8: Foreign-twin retrieval diagnostics. The figure reports whether multilingual publication variants are found and how deeply a model must search before reaching the first cross-language twin.




Figure 9: Embedding separability for same-language and cross-language gold documents. Lower cross-language separability indicates that the retrieval failure is partly representation-level, not only a top- k ranking artifact.




Figure 10: Cross-lingual ranking consistency for the same chemical concept asked in multiple languages. Low overlap means language variants of the same query can return substantially different ranked lists.




Figure 11: Joint language and chemistry failure modes among confused cases. This analysis links the paper’s two central risks: same-language preference and chemically plausible but wrong neighbors.




Figure 12: Question-type effects in the alias-graph benchmark. Some query types are substantially more vulnerable to chemistry-confusability than others.




Figure 13: Confusion as separability collapse. Confused examples have weaker embedding separation between gold and non-gold documents.




Figure 14: Universal attractors in the alias graph: chemically plausible look-alikes that repeatedly steal top rank across models or languages.




Figure 15: Human-written versus machine-translated queries. This supports the paper’s practical separation between generating query variants and relying on human-translated source evidence.




Figure 16: Re-ranker recoverability curves. These curves show how much cross-language evidence is recoverable at different candidate-pool depths and why some failures must be fixed before re-ranking.