

Measuring Cross-Lingual Robustness for Chemistry-Patent Retrieval: Patent-Grounded Multilingual Benchmarks, New Metrics, and a Deployment Decision

Anonymous ACL submission

Abstract

Chemical knowledge is often searched and reused across language boundaries: a user may ask a technical question in one language while the strongest supporting document is written in another. Retrieval systems should therefore be tested on whether they find that cross-language evidence, not only on whether they retrieve an easier same-language match. Yet multilingual chemistry retrieval has little clean public evaluation data, so models are often judged by a single aggregate recall score that can hide poor cross-language retrieval. We address this gap with a benchmark suite containing two public multilingual chemistry QAC retrieval releases, one derived from Google Patents and one from EPO patent data, together with a generation pipeline for cross-lingual retrieval evaluation. These two patent-derived releases provide controlled multilingual technical evidence, while chemistry ontologies enable an auxiliary chemically plausible hard-negative diagnostic. We export the benchmark suite in retrieval-ready form with corpus, queries, qrels, QAC records, and language metadata, enabling evaluation by query language, document language, same-language versus cross-language relevance, and chemical confusability. Aggregate Recall@10 overstates deployment readiness: the best cross-language Recall@10 is about 0.55, same-language advantage remains visible in every query language, and chemically confusable documents outrank gold evidence on 14–48% of publication-lens queries. The benchmark suite and pipeline provide a practical framework for building and auditing specialized multilingual retrieval evaluations in high-trust technical domains. Data and code are available on [Hugging Face \(Google Patents\)](#), [Hugging Face \(EPO\)](#), and [GitHub](#).

1 Introduction

Technical search in chemistry rarely respects a single language boundary. A researcher may ask a question in English, German, French, Spanish, or

another working language, while the strongest evidence may be documented elsewhere: in a translated patent family, in a regional filing, or in technical text whose terminology follows a different linguistic convention. This is not only a translation problem. Chemistry search depends on aliases, abbreviations, formulae, compound families, process descriptions, and near-neighbor concepts that can be semantically close while still being wrong for the user’s question. For a retrieval system used in a high-trust technical workflow, returning a plausible but chemically different document is not a harmless ranking error.

This matters in multinational industrial settings, where technical documents may be stored in several languages but downstream users and systems need the relevant evidence regardless of the query language or document language. The same retrieval requirement appears in RAG systems, agentic workflows, and other decision-support pipelines that depend on finding the right evidence before reasoning over it. Translating every query or document can help in some cases, but it becomes less attractive as language coverage grows because translation adds latency, cost, and another source of technical-term errors.

Dense retrieval with neural sentence and passage embeddings (Reimers and Gurevych, 2019; Karpukhin et al., 2020) is attractive for this setting because multilingual embedding models promise a single retrieval index over many languages (Feng et al., 2022; Chen et al., 2024). In practice, however, the evaluation signal used to choose such a model is often too coarse. A model can achieve a strong average Recall@k by retrieving easy same-language evidence, while failing when the intended evidence is in another language, a bias also observed in recent cross-lingual retrieval analyses (Goworek et al., 2025; Amiraz et al., 2025). The problem is not only whether a multilingual encoder retrieves a relevant document, but whether it

retrieves the right document when query language, document language, and chemical similarity pull the ranking in different directions. This is the failure mode that matters in deployment, and it is precisely the one that an aggregate retrieval score can hide.

Existing multilingual retrieval benchmarks provide important coverage for general cross-lingual search (Zhang et al., 2023; Lawrie et al., 2023; Muennighoff et al., 2023; Enevoldsen et al., 2025), but they do not directly capture this chemistry workflow. We need evaluation data where language variants are controlled, questions are grounded in technical contexts, and hard negatives are chemically plausible rather than arbitrary. Without this structure, an evaluation may reward the model that looks best on average while missing the model behavior that creates practical risk.

We present a multilingual chemistry question-answer-context (QAC) benchmark suite and generation pipeline for cross-lingual retrieval evaluation. Our goal is to evaluate multilingual chemistry retrieval, not patent search alone. We use patents as the benchmark substrate because patent families often contain related technical documents in multiple languages and include chemistry-rich descriptions at scale, building on the long history of patent retrieval evaluation (Piroi et al., 2011).

This paper is therefore an evaluation-framework paper rather than a model leaderboard. Our contributions are: (1) two public multilingual chemistry-patent retrieval benchmark releases, derived from Google Patents and EPO patent data, with query/document language metadata and retrieval-ready corpus, queries, and qrels; (2) a reproducible LLM-assisted QAC generation and human-audit pipeline; (3) a language-aware retrieval evaluation that separates same-language and cross-language evidence; (4) an auxiliary chemistry-confusability stress test based on alias and ontology neighbors; and (5) a deployment-oriented model-selection analysis showing why aggregate recall is not a sufficient deployment dashboard.

2 Related Work

General embedding and multilingual retrieval benchmarks, including MTEB, MMTEB, MIRACL, NeuCLIR, and CLIRMatrix, enable systematic comparison of retrieval and embedding models across languages (Zhang et al., 2023; Lawrie et al., 2023; Muennighoff et al., 2023; Enevold-

sen et al., 2025; Sun and Duh, 2020). Chemistry-specific resources such as ChemTEB, ChemLitQA, ChemComp, ChemKGMultiHopQA, and ChEmbed extend evaluation to chemical embeddings, RAG, and scientific reasoning (Kasmaee et al., 2024; Wellawatte et al., 2024; Khodadad et al., 2026; Astaraki et al., 2026; Kasmaee et al., 2025). However, these lines of work do not jointly provide multilingual chemistry retrieval data with controlled technical variants, explicit same-language and cross-language relevance judgments, and chemically plausible hard negatives. This missing combination motivates the benchmark suite introduced in this paper.

Recent work on cross-lingual retrieval and multilingual RAG shows that retrieval behavior can depend strongly on language alignment, translation effects, and same-language preference (Goworek et al., 2025; Amiraz et al., 2025; Li et al., 2024; Liu et al., 2025; Ki et al., 2025). These findings expose a gap for multilingual chemistry retrieval: aggregate recall alone does not show whether a model retrieves the right evidence across languages or mainly benefits from easier same-language matches. Patent publication variants are useful for studying this gap because they provide comparable multilingual technical evidence under tighter content control than independently collected documents across languages.

Patent retrieval has a long evaluation history, including NTCIR and CLEF-IP (Fujii and Ishikawa, 2002; Piroi et al., 2011), while recent work has revisited patent representation learning and embedding evaluation (Ghosh et al., 2024; Yousefiramandi and Cooney, 2026). Because patent publications often contain related technical disclosures across languages, they can serve as a source of comparable multilingual technical evidence beyond prior-art search alone. Chemical ontology and alias resources, including ChEBI (Hastings et al., 2016), also make it possible to define hard negatives that are plausible in chemistry search but incorrect for a target query. Such chemistry-aware retrieval tests complement ontology-grounded extraction work such as CEAR (Langer et al., 2024): CEAR extracts chemical entities and roles, whereas retrieval evaluation asks whether models distinguish intended evidence from chemically related but incorrect documents.

3 QAC Generation Pipeline

Our method turns multilingual chemistry technical text into auditable question-answer-context (QAC) records for retrieval evaluation. The design goal is not simply to generate fluent questions. Each generated item must identify the evidence document it came from, preserve language metadata, and carry quality signals that make the benchmark suite inspectable after export. Figure 1 summarizes the automated generation and retrieval-evaluation flow.

Corpus and context construction. The input is a chemistry-relevant multilingual corpus. In this work, we use two patent-derived sources, Google Patents and EPO patent data, because patents provide technical descriptions and multilingual publication variants at scale. Both sources are normalized into the same document-language representation before QAC generation. For each document-language pair, the context builder assembles a compact evidence block from the title and abstract, adding first-claim text when available; longer description fields are source-dependent and are not assumed to be present for every document. This step keeps the generation grounded in a specific document rather than asking the model to produce generic chemistry questions from memory.

Candidate generation. For each selected document, language, and generation mode, an LLM produces multiple QAC candidates. Each candidate contains a user-facing question, a short answer, the supporting context, and metadata linking the item back to the source document and language. We use both technical and semantic question modes so that the benchmark suite contains queries that reflect different retrieval intents: some ask for concrete technical details, while others ask for a higher-level description of what the document is about.

Scoring and filtering. Generated candidates are then passed through an LLM-based scorer. The scorer scores whether the answer is faithful to the provided context, whether the question is answerable from that context, whether the language is appropriate, and whether the item is useful as a retrieval query. We retain the best-scored items together with their audit metadata rather than treating generation as a black box. This gives each QAC row a provenance trail: source document, language, generation mode, scores, and scorer feedback.

Human-audited validation sample and export.

Because the scorer is itself model-based, we audit its behavior after generation against a 97-item human-annotated sample from the EPO-derived QAC generation data. This audit is not an automatic stage for every item in Figure 1; instead, it checks whether the automatic grading is directionally reliable enough for benchmark construction. The filtered QAC records are then exported with the retrieval artifacts used in later sections: a corpus table, query table, relevance judgments, and the full QAC table. These artifacts let us evaluate models with standard dense-retrieval tooling while still slicing results by query language, document language, and source provenance.

4 Benchmark Construction

The filtered QAC records are exported as retrieval benchmarks rather than as a single flat table. We release two patent-derived benchmark releases: one based on Google Patents and one based on EPO patent data. Together, these releases form the benchmark suite used for the main retrieval evaluation. Each release contains four linked views: corpus stores the searchable documents and their languages; queries stores the generated questions; qrels stores relevance judgments from queries to documents; and qac preserves the original question, answer, context, source document, generation metadata, and quality scores. This separation lets standard retrieval evaluators consume the data while keeping the provenance needed to audit failures by language and source. Table 1 summarizes the released dataset statistics.

Cross-language relevance. For the multilingual chemistry retrieval benchmark suite, each query is tied to a source technical document. Relevant documents include the source document and its available multilingual variants, so a query can be evaluated against both same-language and cross-language evidence. We keep query language, document language, and source identifiers as explicit metadata. This makes it possible to separate an easy same-language hit from a harder cross-language hit, and to construct “no-home” cases in which the query language has no same-language gold document. These cases are central to the benchmark suite because they remove the shortcut that often inflates aggregate recall.




Figure 1: Overview of the multilingual chemistry retrieval benchmark pipeline. The generation stage samples a source document, builds a grounded title/abstract context, generates candidate queries, scores them, and exports the best-scored items as a retrieval dataset. The evaluation stage embeds queries and documents, ranks results, and reports language-aware, chemistry-confusability, and deployment-oriented diagnostics.

Figure note: Icons adapted from Font Awesome Free, licensed under CC BY 4.0.

Source	Corpus	Queries	Qrels	Cross-lang qrels	Q tokens	Q vocab	C tokens	C vocab
Google Patents	23,787	524	1,284	1,023	11.3 ± 6.3	2,742	145.0 ± 63.8	83,333
EPO	11,315	198	594	396	14.1 ± 5.4	1,371	178.2 ± 71.1	72,256

Table 1: Statistics for the two released patent-derived multilingual QAC retrieval benchmark releases. Question and corpus token columns report mean $\pm$ standard deviation using regex token counts. Vocabularies are unique case-folded tokens.

Chemistry-aware hard negatives. Language is only one source of retrieval error in chemistry. A model may also rank a document about a related compound, synonym, parent class, or neighboring concept above the intended evidence. In addition to the two QAC benchmark releases, we report an auxiliary alias-graph stress test built from chemistry ontology and alias structure. ChEBI concepts and multilingual aliases define target chemical concepts; documents associated with the target concept form the gold set; and taxonomic or alias-graph neighbors provide chemically plausible hard negatives. These negatives are deliberately more informative than random non-relevant documents: they ask whether the retriever can distinguish the right chemical evidence from evidence that is close enough to be tempting.

The resulting evaluation labels each retrieved document by its role in the workflow: same-

language relevant evidence, cross-language relevant evidence, unrelated distraction, or chemically confusable hard negative. The next section uses these labels to evaluate multilingual embedding models beyond aggregate recall.

5 Evaluation Setup

We evaluate multilingual embedding models as drop-in dense retrievers. Each model embeds the query and all corpus documents; documents are ranked by embedding similarity; and retrieval quality is computed from the benchmark suite’s qrels. This setup matches the practical decision faced by a technical search team: choose one multilingual encoder for a shared index, then ask whether it retrieves the right evidence across languages.

Models and protocol. We evaluate nine multilingual or domain-relevant embedding models:

embeddinggemma, bge-m3 (Chen et al., 2024), granite-278m, nomic-v2-moe, qwen3-0.6B, LaBSE (Feng et al., 2022), SapBERT (Liu et al., 2021), e5-large-instruct, and gte-base. All models retrieve against the same multilingual chemistry corpus, and all queries are evaluated with the same relevance judgments. We use standard Recall@k as the baseline retrieval metric because it is familiar and easy to compare, but we treat it as only the starting point: average recall does not say where the relevant document came from or which failure mode produced a miss.

Language-aware views. For each model, we split results by query language, document language, and whether a retrieved relevant document is in the same language as the query or in another language. This gives two central views: same-language retrieval, where a model can often exploit lexical or embedding alignment advantages, and cross-language retrieval, where the model must bridge the language boundary. We also inspect language-pair behavior as a directed matrix, since retrieving a German document from an English query need not be as easy as retrieving an English document from a German query.

Failure analysis. Beyond relevant-document recall, we measure the kinds of mistakes that matter for the workflow. First, we track same-language distractions: cases where a model retrieves a non-relevant document in the query language ahead of relevant cross-language evidence. Second, we use the auxiliary alias-graph stress test to measure chemistry confusion, where a document about a chemically related but incorrect concept outranks the intended evidence. These analyses turn retrieval output into diagnostic categories rather than a single score.

6 Results

We organize the results around the questions the benchmark suite is meant to answer: how models perform across releases and languages, where aggregate recall hides cross-language weakness, whether generated QACs are usable, and whether chemistry-aware hard negatives surface failures that ordinary relevance metrics hide.

Model performance varies by release and language. Figure 2 gives the first-order retrieval picture across all evaluated models and query languages. The strongest models are stable across the

two releases: embeddinggemma has the highest balanced Recall@10 on both Google Patents (0.57) and EPO (0.58), while bge-m3 is close behind, especially on EPO (0.55). The heatmaps also show why a single leaderboard number is incomplete: weaker models degrade unevenly across German, Spanish, French, and Chinese queries, and the EPO release has a different language mix from Google Patents.

Average recall hides cross-language weakness. Figures 3 and 4 show that recall depends on question type, language, and relevance position. Technical questions are harder than semantic questions, and same-language gold evidence is retrieved more easily than cross-language gold evidence in every query language; in the updated Google Patents results, the same-vs-cross gap is roughly 0.21–0.28 Recall@10. A model can therefore look acceptable under an aggregate score while failing on the deployment-relevant case where the evidence is available only in another language.

QAC quality is high enough for benchmark construction. The 97-item human-audited sample from the EPO-derived QAC generation data gives a mean score of 8.33/10, with 94 items falling in the “good” bucket. Technical questions score slightly higher than semantic questions on average (8.72 vs. 7.91), but both modes remain usable for retrieval evaluation. The LLM scorer is slightly stricter than the human annotator (79.0% vs. 83.3%), supporting its use for filtering while still requiring human audit for benchmark-quality claims.

The errors are not only multilingual; they are chemical. The auxiliary alias-graph stress test shows a second failure mode. Even when a retrieved document is topically close, it may be chemically wrong. Figure 5 measures how often a chemically confusable document outranks all gold evidence; in the updated alias-graph results, this occurs on 14–48% of queries depending on the model. These failures are especially important for chemistry search because a nearby compound, parent class, or sibling concept can look plausible while answering the wrong question.

Together, the results support the central lesson of the benchmark suite: multilingual chemistry retrieval should be evaluated with language position and chemical confusability visible, not only with a single averaged retrieval score.




Figure 2: Recall@10 by model and query language for the two benchmark releases. The balanced-mean column summarizes broad model quality, while the language columns show that model rankings and absolute recall vary by query language and source.




Figure 3: Semantic-minus-technical Recall@10 gap by model and language. Nearly every cell is positive, showing that technical chemistry questions are harder to retrieve than semantic questions across both patent sources.

7 Deployment Decision

The evaluation is meant to support a practical model-selection decision: choose one multilingual embedding model for a shared chemistry retrieval index. If the team looked only at aggregate retrieval scores, the decision would be hard to audit because same-language hits and chemically plausible distractors are mixed into one number. Our benchmark suite changes the decision criterion from “which model has the best average Recall@10?” to “which model remains usable when the evidence is cross-language and chemically confusable?”

Capability vs. reading cost. Figure 6 decomposes the deployment cost of cross-lingual retrieval: how deeply a user must read, how much a re-ranker can recover from a candidate pool, and

which failures remain beyond re-ranking. This turns model selection from an average-recall ranking into an operational question about capability, reading depth, and recoverability.

Figure 7 then shows the resulting cost-vs-capability frontier. We use CLIR@10, our cross-language Recall@10 metric, as the capability axis. We use XRC50, the median multiplier in reading depth needed to reach cross-language evidence compared with same-language evidence, as the cost axis. The Pareto frontier has three models: embeddinggemma, bge-m3, and granite-278m. embeddinggemma is the maximum-CLIR corner (about CLIR@10 = 0.55, XRC50 = 5.0), so it is the capability-first choice. If an operating threshold around CLIR@10 = 0.45 is acceptable, bge-m3 becomes the cheaper-to-read admitted model (XRC50 about 3.5). At a lower thresh-




Figure 4: Recall@10 separates into same-language gold retrieval and cross-language gold retrieval. Same-language evidence is consistently easier to retrieve, which explains why average recall can hide the failure most relevant to multilingual chemistry search.




Figure 5: Confusion rate by model and query language: how often a chemically plausible but wrong look-alike outranks every gold document. Lower rates are better, and the spread shows why chemistry-aware hard negatives are needed.

old, granite-278m enters as an even cheaper frontier point (XRC50 about 3.0, CLIR@10 about 0.39). Thus the benchmark suite does not produce a context-free winner; it exposes the trade-off a deployment team must choose.

Decision rule. For a deployment that prioritizes maximum cross-language recall, we would select embeddinggemma. For a deployment with a stricter reading-cost budget and an acceptable CLIR@10 threshold around 0.45, bge-m3 is the more efficient alternative. If a lower recall threshold is acceptable, granite-278m is cheaper still. Figure 6 adds the workflow view: some misses are recoverable with a deeper candidate pool or re-ranker, while the residual floor requires a better first-stage retriever. In all cases, the decision uses language-aware and chemistry-aware diagnostics rather than aggregate recall alone. This is the practical lesson: model selection for high-trust multilingual chem-

istry search should report capability, language robustness, and chemistry-confusability side by side before a retriever is shipped.

Limitations

This benchmark suite is intentionally specialized. Chemistry patents provide a useful substrate for controlled multilingual technical evidence, but results may not transfer directly to other scientific domains, regulatory corpora, or general web retrieval. Patent publication variants also give stronger content control than unrelated multilingual corpora, but they do not guarantee full claim-level equivalence across every language version. The design should therefore be read as a reusable evaluation pattern, not as proof that the same model ordering will hold everywhere or that all patent variants are legally interchangeable.

The QAC pipeline relies on LLM generation and LLM-based verification. The 97-item human-audited sample supports the usefulness of the generated data, but it is limited and does not replace expert patent or chemistry review for every item. Future versions should expand expert annotation, especially for borderline chemical relevance and hard-negative severity.

Finally, the current paper focuses on multilingual retrieval in the two QAC benchmark releases and on auxiliary alias-graph confusability diagnostics. Code-switching and noisy-query infrastructure exist in the project, but we do not claim complete results for those settings here.

8 Conclusion

Multilingual chemistry retrieval needs evaluation data that makes the right failure modes visible. In this paper, we presented a QAC generation and human-audit pipeline for building such data, released two patent-derived benchmark releases in retrieval-ready form, and used the resulting benchmark suite to evaluate dense multilingual retrievers beyond aggregate recall. Language metadata separates same-language from cross-language relevance; QAC provenance makes generated queries auditable; the human-audited sample checks the LLM-based verifier; and the auxiliary alias-graph stress test exposes chemistry-specific confusions.

The central lesson is simple: a model that looks effective on average may still prefer same-language shortcuts or chemically plausible but wrong documents over the intended cross-

Deployment cost of cross-lingual retrieval: how deep to read, how much a re-ranker recovers, and what only alignment can fix




Figure 6: Deployment cost of cross-lingual retrieval. The panels separate how deep a user must read, how much a re-ranker can recover from a candidate pool, and which part of the cross-lingual shortfall remains an alignment-only failure for the first-stage retriever.




Figure 7: Deployment-oriented cost-vs-capability frontier. Cross-language Recall@10 measures capability, while XRC50 measures how much deeper a user must read to reach cross-language evidence. The frontier makes the capability/cost trade-off explicit rather than choosing a model from aggregate recall alone.

language evidence. For high-trust multilingual retrieval, the question is not only whether the system retrieves something relevant on average, but whether it finds the intended evidence when language and chemical similarity create tempting shortcuts.

References

- Chen Amiraz, Yaroslav Fyodorov, Elad Haramaty, Zohar Karnin, and Liane Lewin-Eytan. 2025. The cross-lingual cost: Retrieval biases in RAG over Arabic-English corpora. *arXiv preprint arXiv:2507.07543*.
- Mahdi Astaraki, Mohammad Arshi Saloot, Ali Shiraee Kasmaee, Hamidreza Mahyar, and Soheila Samiee. 2026. When iterative RAG beats ideal evidence: A diagnostic study in scientific multi-hop question answering. *arXiv preprint arXiv:2601.19827*.
- Jianlv Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. 2024. BGE M3-Embedding: Multi-lingual, multi-functionality, multi-granularity text embeddings through self-knowledge distillation. *arXiv preprint arXiv:2402.03216*.
- Kenneth Enevoldsen and 1 others. 2025. MMTEB: Massive multilingual text embedding benchmark. *arXiv preprint arXiv:2502.13595*.
- Fangxiaoyu Feng, Yinfei Yang, Daniel Cer, Naveen Arivazhagan, and Wei Wang. 2022. Language-agnostic BERT sentence embedding. *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 878–891.
- Atsushi Fujii and Tetsuya Ishikawa. 2002. NTCIR-3 patent retrieval experiments at ULIS. *arXiv preprint cs/0206035*.
- Mainak Ghosh, Sebastian Erhardt, Michael E. Rose, Erik Buunk, and Dietmar Harhoff. 2024. PaECTER: Patent-level representation learning using citation-informed transformers. *arXiv preprint arXiv:2402.19411*.
- Roksana Goworek, Olivia Macmillan-Scott, and Eda B. Özyiğit. 2025. What drives cross-lingual ranking? retrieval approaches with multilingual language models. *arXiv preprint arXiv:2511.19324*.
- Janna Hastings, Gareth Owen, Adriano Dekker, Marcus Ennis, Namrata Kale, Venkatesh Muthukrishnan, Steve Turner, Neil Swainston, Pedro Mendes, and Christoph Steinbeck. 2016. ChEBI in 2016: Improved services and an expanding collection of metabolites. *Nucleic Acids Research*, 44(D1):D1214–D1219.
- Vladimir Karpukhin, Barlas Öğuz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen tau Yih. 2020. Dense passage retrieval for open-domain question answering. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 6769–6781.
- Ali Shiraee Kasmaee, Mohammad Khodadad, Mahdi Astaraki, Mohammad Arshi Saloot, Nicholas Sherck, Hamidreza Mahyar, and Soheila Samiee. 2025. ChEmbed: Enhancing chemical literature search through domain-specific text embeddings. *arXiv preprint arXiv:2508.01643*.
- Ali Shiraee Kasmaee, Mohammad Khodadad, Mohammad Arshi Saloot, Nicholas Sherck, Stephen Dokas, Hamidreza Mahyar, and Soheila Samiee. 2024. ChemTEB: Chemical text embedding benchmark, an overview of embedding models performance and efficiency on a specific domain. In *Proceedings of The 4th NeurIPS Efficient Natural Language and Speech Processing Workshop*, volume 262 of *Proceedings of Machine Learning Research*, pages 512–531.
- Mohammad Khodadad, Ali Shiraee Kasmaee, Mahdi Astaraki, Nicholas Sherck, Hamidreza Mahyar, and Soheila Samiee. 2026. Evaluating multi-hop reasoning in large language models: A chemistry-centric benchmark. In *Findings of the Association for Computational Linguistics: EACL 2026*, pages 6117–6143.
- Dayeon Ki, Marine Carpuat, Paul McNamee, Daniel Khashabi, Eugene Yang, Dawn Lawrie, and Kevin Duh. 2025. Linguistic nepotism: Trading-off quality for language preference in multilingual RAG. *arXiv preprint arXiv:2509.13930*.
- Stefan Langer, Fabian Neuhaus, and Andreas Nürnberger. 2024. CEAR: Automatic construction of a knowledge graph of chemical entities and roles from scientific literature. *arXiv preprint arXiv:2407.21708*.
- Dawn Lawrie, Sean MacAvaney, James Mayfield, Paul McNamee, Douglas W. Oard, Luca Soldaini, and Eugene Yang. 2023. Overview of the TREC 2022 NeuCLIR track. In *Proceedings of the Thirty-First Text REtrieval Conference (TREC)*.
- Bryan Li, Fiona Luo, Samar Haider, Adwait Agashe, Tammy Li, Runqi Liu, Muqing Miao, Shriya Ramakrishnan, Yuan Yuan, and Chris Callison-Burch. 2024. BordIRlines: A dataset for evaluating cross-lingual retrieval-augmented generation. *arXiv preprint arXiv:2410.01171*.
- Fangyu Liu, Ehsan Shareghi, Zaiqiao Meng, Marco Basaldella, and Nigel Collier. 2021. Self-alignment pretraining for biomedical entity representations. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pages 4228–4238.
- Wei Liu, Sony Trenous, Leonardo F. R. Ribeiro, Bill Byrne, and Felix Hieber. 2025. XRAG: Cross-lingual retrieval-augmented generation. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 15669–15690.
- Niklas Muennighoff, Nouamane Tazi, Loïc Magne, and Nils Reimers. 2023. MTEB: Massive text embedding benchmark. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, pages 2014–2037.
- Florina Piroi, Mihai Lupu, Allan Hanbury, and Veronika Zenz. 2011. CLEF-IP 2011: Retrieval

in the intellectual property domain. In *CLEF (Notebook Papers/Labs/Workshop)*. Multilingual (en/de/fr) patent prior-art retrieval; see also CLEF-IP 2009–2013 overviews.

Nils Reimers and Iryna Gurevych. 2019. [Sentence-BERT: Sentence embeddings using Siamese BERT-networks](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3982–3992.

Shuo Sun and Kevin Duh. 2020. [CLIRMatrix: A massively large collection of bilingual and multilingual datasets for cross-lingual information retrieval](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 4160–4170.

Geemi P. Wellawatte, Huixuan Guo, Magdalena Lederbauer, Anna Borisova, Matthew Hart, Marta Brucka, and Philippe Schwaller. 2024. [ChemLit-QA: A human evaluated dataset for chemistry RAG tasks](#). In *NeurIPS 2024 Workshop on Machine Learning and the Physical Sciences*.

Amirhossein Yousefiramandi and Ciaran Cooney. 2026. [Benchmarking patent embeddings: A multi-task evaluation of 22 models across retrieval, classification, and clustering](#). *arXiv preprint arXiv:2605.24297*.

Xinyu Zhang, Nandan Thakur, Odunayo Ogundepo, Ehsan Kamalloo, David Alfonso-Hermelo, Xiaoguang Li, Qun Liu, Mehdi Rezagholizadeh, and Jimmy Lin. 2023. [MIRACL: A multilingual retrieval dataset covering 18 diverse languages](#). *Transactions of the Association for Computational Linguistics*, 11:1114–1131.




Figure 8: Same-language and cross-language gold retrieval by query language. The persistent gap explains why aggregate recall can overstate deployment readiness.




Figure 9: Recall@10 by query language, pooled across models with per-model points. This supporting view shows that aggregate retrieval quality varies by query language before separating same-language and cross-language relevance.

A Appendix

This appendix collects supporting analyses that are useful for interpretation but too detailed for the six-page body.

B Additional Cross-Lingual Retrieval Diagnostics

These diagnostics expand the paper’s main cross-language retrieval result. They show that the same-language advantage is not a single-score artifact: it appears as query-language variation, same-language gold advantage, transfer across patent offices, and weaker score separability when gold evidence competes with look-alike documents.

C Additional Alias-Graph Diagnostics

The alias-graph analyses isolate the chemistry side of the retrieval problem. They test whether multilingual queries for the same concept produce con-




Figure 10: Model standing transfers across patent offices. Google Patents and EPO scores are strongly correlated, suggesting that the benchmark-suite signal is not tied to only one source.




Figure 11: Score separability collapses when the model confuses gold evidence with a chemical look-alike. Confused cases approach a coin-flip separation between the right and wrong document.

sistent rankings and whether chemically plausible neighbors become systematic distractors. Together, these figures explain why relevance in chemistry search cannot be reduced to topical similarity alone.

D Deployment-Oriented Supporting Analyses

These figures support the deployment interpretation in the main paper. They are secondary to the main cost-vs-capability decision, but they show two practical constraints: query translation can change retrieval behavior, and re-rankers can only recover cross-language evidence that appears within a feasible candidate pool.




Figure 12: Cross-lingual ranking consistency for the same chemical concept asked in multiple languages. Low overlap means language variants of the same query can return substantially different ranked lists.




Figure 13: Confusion rate by model and query language. The figure shows how often a chemically plausible look-alike outranks every gold document.

E Question-Type and Metric Robustness

The main text reports that technical and semantic questions are both usable, but the new diagnostics show that they are not equally easy retrieval tasks. Technical questions are consistently harder, and that penalty appears across patent sources, models, languages, and retrieval metrics. This matters because a retrieval evaluation made only from broad semantic questions would overstate readiness for technical chemistry search.

F Cross-Source and Language-Denominator Checks

The paper now uses both Google Patents and EPO-derived benchmark releases. These additional figures check whether the main model ordering and difficulty patterns transfer across the two sources, and they document the denominators behind language-balanced reporting.




Figure 14: Structure-style questions are the hardest alias-graph cases: they combine low recall with high chemistry-confusion.




Figure 15: Confusion as separability collapse. Confused examples have weaker score separation between gold and look-alike documents.

G Additional Deployment-Budget Diagnostics

These deployment views unpack the cost-vs-capability frontier in the main paper. They separate reading cost, re-ranker recoverability, and the residual alignment-only failure floor, giving practitioners a more concrete way to decide whether a retriever failure can be fixed downstream.




Figure 16: Universal attractors in the alias graph: chemically plausible look-alikes that repeatedly steal top rank across models or languages.




Figure 17: Chemistry-confusability with attractor detail. The left panel shows how often a plausible but wrong document outranks all gold evidence; the right panel identifies which look-alike concepts most often cause the failure.




Figure 18: Human-original versus machine-translated questions on Google Patents. The comparison checks whether the retrieval signal is dominated by translation artifacts.




Figure 19: Re-ranker recoverability curves. The curves show how much cross-language evidence is recoverable at different candidate-pool depths and mark the knee point for each model.




Figure 20: Technical questions retrieve worse than semantic questions across models and patent offices. The gap is visible for both Google Patents and EPO, so the question-mode distinction is not a single-source artifact.




Figure 21: The technical-question penalty is a distribution shift, not only a few outlier queries. Semantic questions place more mass at higher Recall@10, and the per-model view shows the same ordering across both sources.




Figure 22: Semantic-minus-technical Recall@10 gaps by model and language. Most cells are positive, showing that technical questions are harder across many language and model combinations.




Figure 23: The technical-question penalty remains under MRR@10 and Hit@10. This checks that the result is not an artifact of choosing Recall@10 as the main retrieval metric.




Figure 24: Within technical questions, parameter/condition and method questions are hardest, while outcome questions are easier. The result helps explain why technical retrieval cannot be summarized by a single question category.




Figure 25: Few rank crossings occur between Google Patents and EPO. The strongest models remain near the top across both patent offices, while weaker models remain lower.




Figure 26: The semantic-over-technical retrieval gap transfers across patent offices. Most models fall in the positive-positive quadrant, meaning semantic questions are easier on both sources.

Per-language denominators are unequal — why per-language means must be language-balanced




Figure 27: Per-language denominators for queries, gold qrels, and haystack documents are unequal. Language-balanced metrics prevent high-resource languages from dominating the summary score.

Deployment cost of cross-lingual retrieval: how deep to read, how much a re-ranker recovers, and what only alignment can fix




Figure 28: Deployment cost decomposed into reading cost, re-ranker recoverability, and residual alignment-only failure. The triptych shows which failures are reachable with a deeper candidate pool and which require a better first-stage retriever.




Figure 29: Cost-vs-capability frontier with an inset re-ranker budget curve. This combines model choice and downstream recoverability in a single deployment-oriented view.

Re-ranker budget: most of the cross-lingual gap is reachable in a top-100 pool, a floor beyond it is not




Figure 30: Re-ranker budget decomposition. Green regions show failures recoverable in top-10 or top-100 pools, while the red region is an alignment-only floor beyond the re-ranker budget.